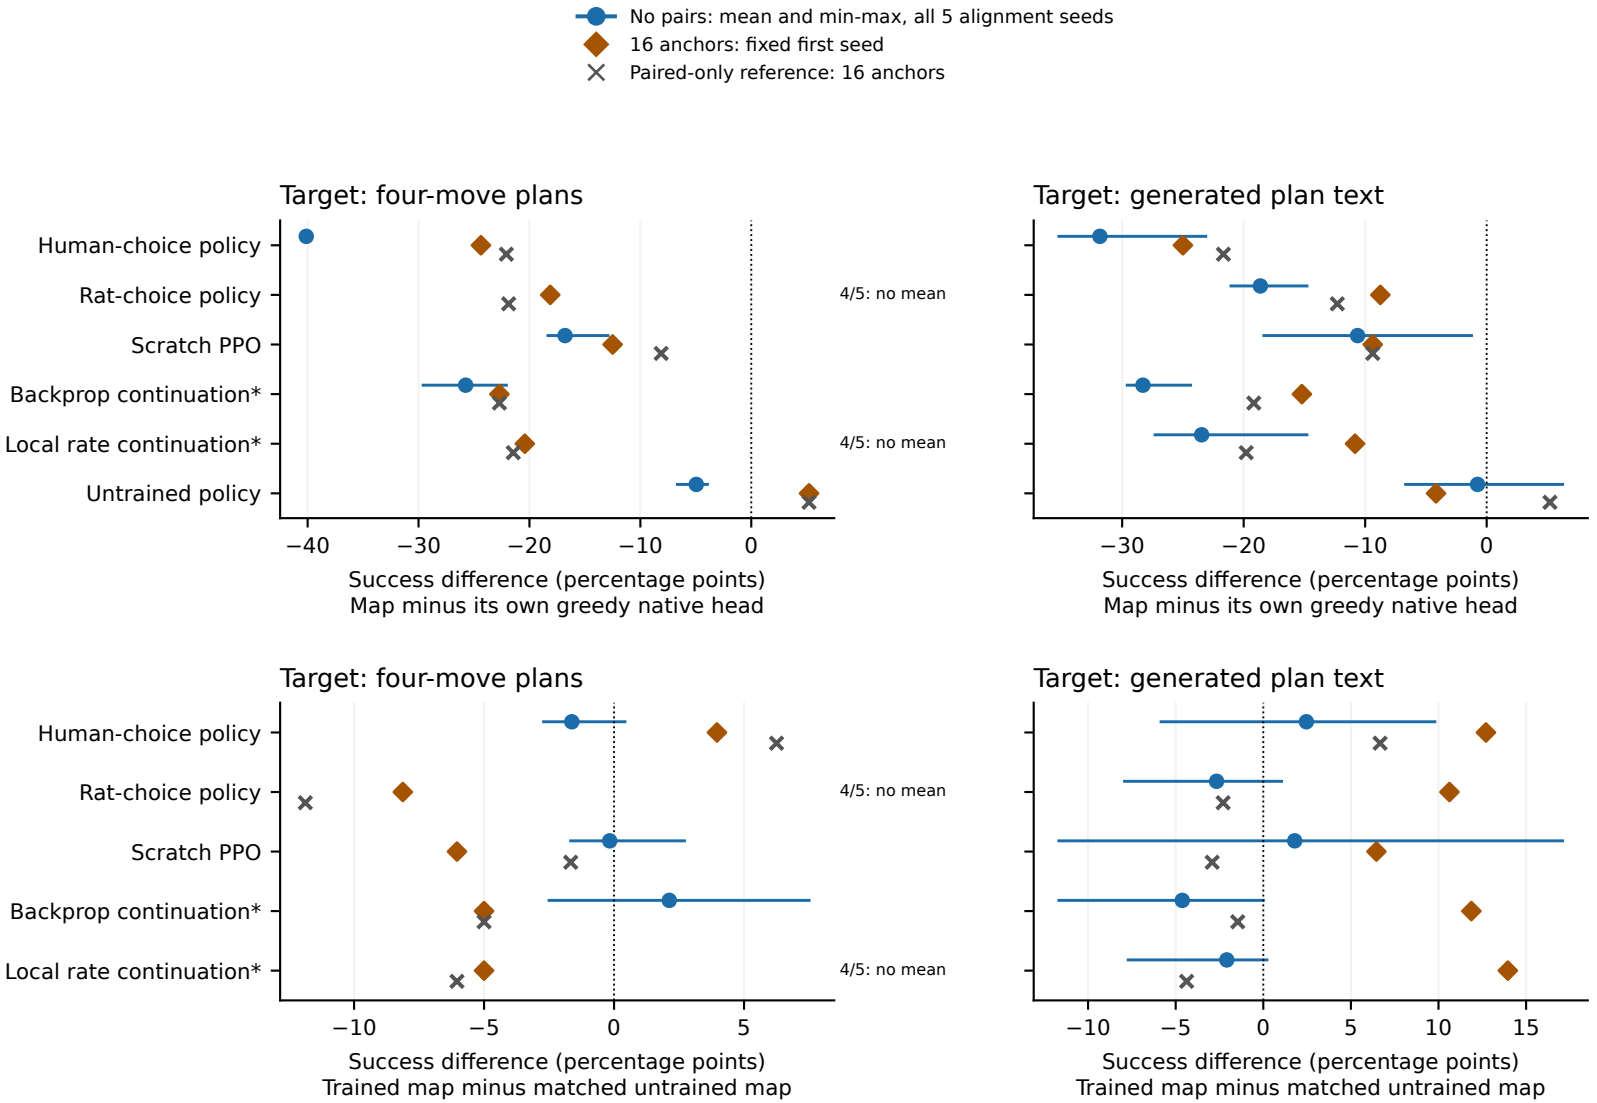

One policy-training seed (60); alignment seeds measure optimizer variation, not training/biological replication. No confidence intervals.
The same 16 exposed rooms; 30 episodes/room. Native head is a comparator, not an upper bound. Sensory-only maps lack matched heads.
*Both continuations start from backprop-trained policies, then receive the same extra source frames; local updates are rate based.
Text is a fixed generated command catalog. Rat/human labels identify supervision sources, not online animal or neural inputs.